

Simple Discretization Methods: Binning

1. Equal-Width (Distance) Partitioning

Formula:

$$W = \frac{B - A}{N}$$

where:

- W = Width of each interval
- A = Minimum value in the dataset
- B = Maximum value in the dataset
- N = Number of bins

Example: Given the dataset:

{1, 3, 7, 10, 15, 18, 22, 25}

Let $N = 3$, $A = 1$, $B = 25$.

Step 1: Compute Interval Width

$$W = \frac{25 - 1}{3} = \frac{24}{3} = 8$$

Step 2: Define Bins

- **Bin 1:** $[1, 9) \rightarrow$ contains {1, 3, 7}
- **Bin 2:** $[9, 17) \rightarrow$ contains {10, 15}
- **Bin 3:** $[17, 25] \rightarrow$ contains {18, 22, 25}

This method is **simple and intuitive**, but **outliers** and **skewed data** can make it less effective.

2. Equal-Depth (Frequency) Partitioning

Concept: Each bin should have **approximately the same number of elements**.

Example: Given the dataset:

{1, 3, 7, 10, 15, 18, 22, 25}

Let $N = 3$. Since there are **8 elements**, each bin should have about $\frac{8}{3} \approx 3$ elements.

Step 1: Sort Data

{1, 3, 7, 10, 15, 18, 22, 25}

Step 2: Assign Bins

- **Bin 1:** {1, 3, 7}
- **Bin 2:** {10, 15, 18}
- **Bin 3:** {22, 25}

This method ensures **equal distribution** of data points in each bin, making it **better for skewed data**.

Bin Processing Methods

1. Bin Means

Concept: Replace all values in a bin with the **mean** (average) of that bin.

Example: Using the same bins from equal-width partitioning:

- **Bin 1:** $\{1, 3, 7\} \rightarrow \text{Mean} = \frac{1+3+7}{3} = 3.67 \rightarrow \{3.67, 3.67, 3.67\}$
- **Bin 2:** $\{10, 15\} \rightarrow \text{Mean} = \frac{10+15}{2} = 12.5 \rightarrow \{12.5, 12.5\}$
- **Bin 3:** $\{18, 22, 25\} \rightarrow \text{Mean} = \frac{18+22+25}{3} = 21.67 \rightarrow \{21.67, 21.67, 21.67\}$

Pros: Reduces variance. **Cons:** May hide important variations in data.

2. Bin Medians

Concept: Replace all values in a bin with the **median** of that bin. The **median** is the middle value when the numbers are sorted. If there is an even number of elements, the median is the average of the two middle values.

Example:

- **Bin 1:** $\{1, 3, 7\} \rightarrow \text{Sorted: } \{1, 3, 7\}, \text{Median} = 3 \rightarrow \{3, 3, 3\}$
- **Bin 2:** $\{10, 15\} \rightarrow \text{Sorted: } \{10, 15\}, \text{Median} = \frac{10+15}{2} = 12.5 \rightarrow \{12.5, 12.5\}$
- **Bin 3:** $\{18, 22, 25\} \rightarrow \text{Sorted: } \{18, 22, 25\}, \text{Median} = 22 \rightarrow \{22, 22, 22\}$

Pros: Handles outliers better than means. **Cons:** May not capture overall distribution well.

3. Bin Boundaries

Concept: Replace each value with the nearest boundary value in the bin.

Example:

- **Bin 1:** $\{1, 3, 7\} \rightarrow \text{Boundaries: } 1, 7 - 1 \rightarrow 1, 3 \rightarrow 1, 7 \rightarrow 7 \rightarrow \{1, 1, 7\}$
- **Bin 2:** $\{10, 15\} \rightarrow \text{Boundaries: } 10, 15 - 10 \rightarrow 10, 15 \rightarrow 15 \rightarrow \{10, 15\}$
- **Bin 3:** $\{18, 22, 25\} \rightarrow \text{Boundaries: } 18, 25 - 18 \rightarrow 18, 22 \rightarrow 25, 25 \rightarrow 25 \rightarrow \{18, 25, 25\}$

Pros: Good for preserving extreme values. **Cons:** Can distort middle-range values.

Data Transformation: Normalization

1. Min-Max Normalization

Formula:

$$v' = \frac{v - \min_A}{\max_A - \min_A} \times (\text{new_max}_A - \text{new_min}_A) + \text{new_min}_A$$

where:

- v = Original value
- \min_A = Minimum value of the original dataset
- \max_A = Maximum value of the original dataset
- new_min_A = Minimum value of the new range
- new_max_A = Maximum value of the new range
- v' = Normalized value

Example: Convert age = 30 to a range [0, 1], where min = 10 and max = 80:

$$v' = \frac{30 - 10}{80 - 10} = \frac{20}{70} = \frac{2}{7} \approx 0.2857$$

This transformation scales the value to fall between the new range, in this case, between 0 and 1.

2. Z-Score Normalization

Formula:

$$v' = \frac{v - \text{mean}_A}{\text{stand_dev}_A}$$

where:

- v = Original value
- mean_A = Mean of the original dataset
- stand_dev_A = Standard deviation of the original dataset
- v' = Normalized value

Rules:

- **Mean rule:** The Z-score normalization shifts the dataset such that the mean of the transformed data becomes 0.
- **Standard deviation rule:** The Z-score normalization scales the dataset so that the standard deviation of the transformed data becomes 1.

The Z-score normalization transforms the data to have a ****mean of 0**** and a ****standard deviation of 1****, making it useful for comparing data from different distributions.

3. Normalization by Decimal Scaling

Formula:

$$v' = \frac{v}{10^j}$$

where j is the smallest integer such that:

$$\text{Max}(|v'|) < 1$$

This transformation scales the data by dividing each value by a power of 10, where the exponent j is chosen to ensure the maximum absolute value of v' is less than 1.